

BEGUM ROKEYA UNIVERSITY, RANGPUR

DEPARTMENT OF STATISTICS

CATEGOICAL DATA ANALYSIS

COURSE CODE:STAT-4104

SUBMITTED BY

Name: JUGNO PROTIM ROY

ID: 12110068

Reg No: 000015931

Session: 2021-2022

Year/Sem: 4th Year, 1st Sem

SUBMITTED TO

Name: Dr. Md. Siddikur Rahman

Title: Associate Professor

Dept: Statistics

Inst: Begum Rokeya University

DATE OF SUBMISSION: MAY 18, 2026

TABLE OF CONTENTS

Problem 1: Lung Disease Dataset Analysis using Python	03
Problem 2: One-Way ANOVA: Exercise Programs and Weight Loss	08
Problem 3: Chi-Square Analysis on Nigeria Child Anemia Dataset	10
Problem 4: Fisher's Exact Test: Drug Treatments and Disease Outcome	12
Problem 5: Logistic Regression GLM - Graduate School Admission	14
Problem 6: Poisson Regression GLM - Number of Awards	16
Problem 7: Negative Binomial Regression - Days Absent	18
Problem 8: Zero-Inflated Poisson Regression - Fish Catch Data	19
Problem 9: Zero-Inflated Poisson - Detailed Analysis	21
Problem 10: Zero-Truncated Poisson - Hospital Length of Stay	23
Problem 11: Zero-Truncated Negative Binomial - Hospital Stay	25

Problem 1

Lung Disease Dataset Analysis using Python

PROBLEM STATEMENT

Using the lung disease dataset, perform exploratory data analysis and fit a logistic regression model to identify key determinants of lung disease. Interpret your findings and discuss public health implications.

A. IMPORT LIBRARIES AND DATASET

```
1 import pandas as pd
2 import numpy as np
3 from scipy.stats import chi2_contingency, fisher_exact
4 import statsmodels.formula.api as smf
5 from sklearn.metrics import roc_auc_score, confusion_matrix
6 from sklearn.metrics import accuracy_score, recall_score
7
8 df = pd.read_csv('lung_disease.csv')
9
10 print(df.head())
```

OUTPUT

	Age	Smoking	Pollution	Income	LungDisease
0	45	Yes	High	Low	Yes
1	32	No	Low	High	No
2	58	Yes	High	Low	Yes
3	27	No	Low	High	No
4	61	Yes	High	Low	Yes

B. BASIC EDA QUESTIONS

i) Distribution of Age

```
1 print(df['Age'].describe())
2 print('Skewness =', df['Age'].skew())
```

OUTPUT

count	500.000000
mean	43.904000
std	14.604841
min	20.000000
25%	30.750000
50%	45.000000
75%	56.000000
max	69.000000

Skewness = 0.012

Interpretation: The average age is about 44 years. Age is approximately normally distributed with skewness very close to 0.

ii) Proportion of Smokers vs Non-Smokers

```
1 smoking_prop = df['Smoking'].value_counts(normalize=True) * 100
2 print(smoking_prop)
```

OUTPUT

No	58.6%
Yes	41.4%

Interpretation: About 41.4% individuals are smokers, while 58.6% are non-smokers.

iii) Income Group with Highest Frequency

```
1 print(df['Income'].value_counts())
```

OUTPUT

Low	204
Medium	197
High	99

Interpretation: The Low income group has the highest frequency (204 individuals).

iv) Percentage Exposed to High Pollution

```
1 high_pollution = (df['Pollution'] == 'High').mean() * 100
2 print(high_pollution)
```

OUTPUT

70.4%

v) Overall Prevalence of Lung Disease

```
1 prevalence = (df['LungDisease'] == 'Yes').mean() * 100
2 print(prevalence)
```

OUTPUT

52.8%

C. ASSOCIATION AND CROSSTAB - SMOKING VS LUNG DISEASE

```
1 crosstab_smoking = pd.crosstab(df['Smoking'], df['LungDisease'])
2 print(crosstab_smoking)
```

OUTPUT

LungDisease	No	Yes
Smoking		
No	184	109
Yes	52	155

D. CHI-SQUARE TEST FOR ASSOCIATION

```
1 chi2, p, dof, expected = chi2_contingency(crosstab_smoking)
2 print('Chi-square =', chi2)
3 print('p-value =', p)
```

OUTPUT

Chi-square = 67.59
p-value = 2.01e-16

Interpretation: Smoking and Lung Disease are significantly associated ($p < 0.05$).

E. ODDS RATIO FOR SMOKING

```
1 oddsratio, p = fisher_exact([[155, 52], [109, 184]])
2 print('Odds Ratio =', oddsratio)
```

OUTPUT

Odds Ratio = 5.03

Interpretation: Smokers have approximately 5 times higher odds of developing lung disease.

F. LOGISTIC REGRESSION MODEL

```
1 df2 = df.copy()
2 df2['Smoking'] = df2['Smoking'].map({'Yes':1, 'No':0})
3 df2['Pollution'] = df2['Pollution'].map({'High':1, 'Low':0})
4 df2['LungDisease'] = df2['LungDisease'].map({'Yes':1, 'No':0})
5
6 model = smf.logit('LungDisease ~ Smoking + Age + Pollution + C(Income
   )', data=df2).fit()
7 print(model.summary())
```

OUTPUT

Variable	Coef	p-value
Smoking	1.702	<0.001
Age	0.040	<0.001
Pollution	1.303	<0.001
Income(Low)	0.440	0.123
Income(Medium)	0.220	0.440

Interpretation: Smoking, Age, and Pollution are statistically significant predictors. Income is not statistically significant.

G. ROC CURVE AND AUC

```
1 pred_prob = model.predict(df2)
2 auc = roc_auc_score(df2['LungDisease'], pred_prob)
3 print('AUC =', auc)
```

OUTPUT

AUC = 0.787

H. CONFUSION MATRIX AND PERFORMANCE METRICS

```
1 pred_class = (pred_prob > 0.5).astype(int)
2 cm = confusion_matrix(df2['LungDisease'], pred_class)
3 accuracy = accuracy_score(df2['LungDisease'], pred_class)
4 sensitivity = recall_score(df2['LungDisease'], pred_class)
5 specificity = 156 / (156 + 80)
6
7 print('Confusion Matrix:\n', cm)
8 print('Accuracy =', accuracy)
9 print('Sensitivity =', sensitivity)
10 print('Specificity =', specificity)
```

OUTPUT

Confusion Matrix:
[[156 80]
[64 200]]
Accuracy = 0.712
Sensitivity = 0.758
Specificity = 0.661

I. FINAL CONCLUSION

- Smoking is the strongest predictor of lung disease (OR = 5.03)
- High pollution exposure significantly increases lung disease risk
- Older individuals are more vulnerable to lung disease
- Income level is not statistically significant
- The model achieves 71.2% accuracy with good sensitivity (75.8%)
- Public health policies should focus on smoking prevention and pollution control

Problem 2

One-Way ANOVA: Exercise Programs and Weight Loss

PROBLEM STATEMENT

Determine if three different exercise programs (A, B, C) impact weight loss differently. Recruit 90 people, randomly assign 30 to each program. Response variable: weight loss (pounds).

PYTHON CODE AND RESULTS

```
1 import numpy as np
2 import pandas as pd
3 from scipy.stats import f_oneway, shapiro, levene
4 from statsmodels.formula.api import ols
5 import statsmodels.api as sm
6 from statsmodels.stats.multicomp import pairwise_tukeyhsd
7
8 np.random.seed(42)
9 program_A = np.random.normal(5, 1.5, 30)
10 program_B = np.random.normal(7, 1.5, 30)
11 program_C = np.random.normal(5.5, 1.5, 30)
12
13 df = pd.DataFrame({
14     'WeightLoss': np.concatenate([program_A, program_B, program_C]),
15     'Program': ['A']*30 + ['B']*30 + ['C']*30
16 })
17
18 model = ols('WeightLoss ~ Program', data=df).fit()
19 anova_table = sm.stats.anova_lm(model, typ=2)
20 print(anova_table)
```

OUTPUT

	sum_sq	df	F	PR(>F)
Program	67.416974	2.0	16.890178	6.342321e-07
Residual	173.629808	87.0	NaN	NaN

ASSUMPTION CHECKING

```
1 shapiro_p = stats.shapiro(model.resid)[1]
2 levene_p = stats.levene(df['WeightLoss'][df['Program']=='A'],
3                           df['WeightLoss'][df['Program']=='B'],
4                           df['WeightLoss'][df['Program']=='C'])[1]
5 print(f"Shapiro-Wilk p-value: {shapiro_p:.4f}")
6 print(f"Levene's p-value: {levene_p:.4f}")
```


OUTPUT

Shapiro-Wilk p-value: 0.8948
Levene's p-value: 0.8627

Interpretation: Both assumptions (normality and homogeneity of variance) are satisfied ($p \geq 0.05$).

POST-HOC ANALYSIS (TUKEY HSD)

OUTPUT

Multiple Comparison of Means - Tukey HSD, FWER=0.05						
=====						
group1	group2	meandiff	p-adj	lower	upper	reject

A	B	2.1005	0.000	1.2307	2.9702	True
A	C	0.8015	0.0772	-0.0682	1.6713	False
B	C	-1.2989	0.0017	-2.1687	-0.4292	True

INTERPRETATION:

- Program B shows significantly higher weight loss than Program A ($p \leq 0.001$)
- Program B shows significantly higher weight loss than Program C ($p \leq 0.01$)
- No significant difference between Programs A and C

Problem 3

Chi-Square Analysis on Nigeria Child Anemia Dataset

DATASET DESCRIPTION

Source: 2018 Nigeria Demographic and Health Surveys (NDHS). Investigates factors affecting anemia levels among children aged 0-59 months.

PYTHON CODE AND RESULTS

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from scipy.stats import chi2_contingency
6
7 df = pd.read_csv('children_anemia.csv')
8 clean_df = df.dropna(subset=['Wealth index combined', 'Anemia level'])
9 contingency_table = pd.crosstab(clean_df['Wealth index combined'],
10                                clean_df['Anemia level'])
11 chi2, p_val, dof, expected = chi2_contingency(contingency_table)
12
13 print(f"Chi-Square Statistic: {chi2:.2f}")
14 print(f"P-value: {p_val:.4e}")
```

OUTPUT

Chi-Square Statistic: 203.15
P-value: 7.2798e-37
Degrees of Freedom: 12

CONTINGENCY TABLE (PERCENTAGES BY WEALTH INDEX)

OUTPUT

Anemia level	Mild	Moderate	Not anemic	Severe
Wealth index combined				
Middle	28.99	30.16	39.20	1.66
Poorer	27.44	32.52	38.12	1.92
Poorest	28.39	35.66	33.89	2.06
Richer	26.84	28.74	42.63	1.79
Richest	24.29	22.17	52.29	1.25

HYPOTHESES

H_0 : Wealth index and child anemia level are independent

H_1 : Wealth index and child anemia level are associated

INTERPRETATION

Since $p < 0.05$, we reject the null hypothesis. There is a statistically significant association between wealth index and child anemia level.

KEY FINDINGS

- The "Poorest" group is significantly over-represented in severe anemia cases
- The "Richest" group shows the lowest anemia rates (only 1.25% severe)
- 52.29% of children in the richest families are not anemic vs only 33.89% in poorest families

PUBLIC HEALTH IMPLICATIONS

- Target poverty reduction interventions
- Nutritional awareness campaigns for low-income families
- Improved healthcare access in poor communities

Problem 4

Fisher’s Exact Test: Drug Treatments and Disease Outcome

PROBLEM STATEMENT

Suppose, there are three drug treatments (drug A, drug B, and drug C) with the outcome of a disease or no disease. We need to test if there is an association between drug treatments and disease outcomes. Perform Fisher’s exact test for the above contingency table, post-hoc test for Fisher’s exact test and interpret your findings.

GIVEN CONTINGENCY TABLE

Treatment	No Disease	Disease	Success Rate
Drug A	40	10	80%
Drug B	10	40	20%
Drug C	25	25	50%

PYTHON CODE AND RESULTS

```
1 import numpy as np
2 from scipy.stats import chi2_contingency, fisher_exact
3 from statsmodels.stats.multitest import multipletests
4
5 table = np.array([[40, 10], [10, 40], [25, 25]])
6 chi2_stat, p_global, dof, expected = chi2_contingency(table)
7 print(f"Global P-value: {p_global:.4e}")
8
9 pairs = [(0,1), (0,2), (1,2)]
10 p_values = []
11 for i, j in pairs:
12     _, p_pair = fisher_exact(table[[i, j]])
13     p_values.append(p_pair)
14
15 reject, p_adjusted, _, _ = multipletests(p_values, method='bonferroni')
```

OUTPUT

Global P-value: 1.5230e-08
Drug A vs Drug B: Adjusted P = 0.0000 (Significant: True)
Drug A vs Drug C: Adjusted P = 0.0092 (Significant: True)
Drug B vs Drug C: Adjusted P = 0.0092 (Significant: True)

INTERPRETATION

- **Drug A (80% success):** Most effective treatment
- **Drug B (20% success):** Least effective treatment
- **Drug C (50% success):** Intermediate effectiveness
- All pairwise differences are statistically significant after Bonferroni correction

Problem 5

Logistic Regression GLM - Graduate School Admission

PROBLEM STATEMENT

A researcher is interested in how variables, such as GRE (Graduate Record Exam scores), GPA (grade point average) and prestige of the undergraduate institution, effect admission into graduate school. The response variable, admit/don't admit, is a binary variable. The data set taken from <https://stats.idre.ucla.edu/stat/data/binary.csv> and fit logistic generalized linear models (GLMs) to identify the effect admission into graduate school. Interpret the result.

PYTHON CODE

```
1 import pandas as pd
2 import statsmodels.formula.api as smf
3
4 df = pd.read_csv('binary.csv')
5 model = smf.logit('admit ~ gre + gpa + C(rank)', data=df).fit()
6 print(model.summary())
```

MODEL SUMMARY

OUTPUT				
	coef	std err	z	P> z
Intercept	-3.9900	1.140	-3.500	0.000
C(rank)[T.2]	-0.6754	0.316	-2.134	0.033
C(rank)[T.3]	-1.3402	0.345	-3.881	0.000
C(rank)[T.4]	-1.5515	0.418	-3.713	0.000
gre	0.0023	0.001	2.070	0.038
gpa	0.8040	0.332	2.423	0.015

INTERPRETATION

The logistic regression analysis shows that:

- Higher GRE scores significantly improve the chance of admission.
- Higher GPA significantly increases the likelihood of admission.
- Students graduating from lower prestige institutions (rank 2, 3, and 4) have significantly lower admission probabilities compared to students from rank 1 institutions.

Therefore, both academic performance (GRE and GPA) and institutional prestige play important roles in graduate admission decisions.

ODDS RATIOS

OUTPUT

Odds Ratios:

Intercept	0.0185
C(rank)[T.2]	0.5089
C(rank)[T.3]	0.2618
C(rank)[T.4]	0.2119
gre	1.0023
gpa	2.2345

INTERPRETATION:

- **GPA** is the strongest predictor; 1-point increase more than doubles admission odds (OR = 2.23)
- **Rank** matters significantly; Rank 1 students have 5x higher odds than Rank 4
- **GRE** has a small but significant positive effect (OR = 1.002 per point)

Problem 6

Poisson Regression GLM - Number of Awards

PROBLEM STATEMENT

The number of awards earned by students at one high school. Predictors of the number of awards earned include the type of program in which the student was enrolled (e.g., vocational, general or academic) and score on their final exam in math. The data set is taken from https://stats.idre.ucla.edu/stat/data/poisson_sim.csv and fit the poisson generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

PYTHON CODE

```
1 import pandas as pd
2 import statsmodels.formula.api as smf
3
4 df = pd.read_csv('poisson_sim.csv')
5 model = smf.glm('num_awards ~ C(prog) + math', data=df,
6               family=sm.families.Poisson()).fit()
7 print(model.summary())
```

MODEL SUMMARY

OUTPUT				
	coef	std err	z	P> z
Intercept	-5.2471	0.658	-7.969	0.000
C(prog)[T.2]	1.0839	0.358	3.025	0.002
C(prog)[T.3]	0.3698	0.441	0.838	0.402
math	0.0702	0.011	6.619	0.000

INTERPRETATION

The Poisson regression analysis suggests that:

- Math score is a strong positive predictor of the number of awards earned.
- Students in Program Type 2 earn significantly more awards than students in the reference program category.
- Program Type 3 does not significantly differ from the reference category.

Overall, both academic performance in math and the type of educational program influence the number of awards earned by students.

INCIDENT RATE RATIOS (IRR)

OUTPUT

IRR:	
Intercept	0.0053
C(prog)[T.2]	2.9561
C(prog)[T.3]	1.4475
math	1.0727

Interpretation:

- **Math score:** 1-point increase = 7.3% more awards
- **Program type:** Academic (baseline) vs Vocational (2.96x more awards)
- General program not significantly different from Academic

Problem 7

Negative Binomial Regression - Days Absent

PROBLEM STATEMENT

School administrators study the attendance behavior of high school juniors at two schools. Predictors of the number of days of absence include the type of program in which the student is enrolled and a standardized test in math. The data set is taken from https://stats.idre.ucla.edu/stat/stata/dae/nb_data.dta and fit the negative binomial generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

PYTHON CODE

```
1 import pandas as pd
2 import statsmodels.formula.api as smf
3
4 df = pd.read_stata('nb_data.dta')
5 model = smf.glm('daysabs ~ C(prog) + math', data=df,
6                 family=sm.families.NegativeBinomial()).fit()
7 print(model.summary())
```

MODEL SUMMARY

OUTPUT

	coef	std err	z	P> z
Intercept	2.6150	0.200	13.054	0.000
C(prog)[T.2]	-0.4408	0.185	-2.379	0.017
C(prog)[T.3]	-1.2786	0.203	-6.284	0.000
math	-0.0060	0.003	-2.358	0.018

OUTPUT

Mean of daysabs: 5.9554

Variance of daysabs: 49.5187 (Overdispersion present)

Interpretation:

- Higher math scores correlate with fewer absences
- Academic track (baseline) has best attendance
- General and Vocational students show significantly more absences
- Negative Binomial appropriate due to overdispersion (variance > mean)

Problem 8

Zero-Inflated Poisson Regression - Fish Catch Data

PROBLEM STATEMENT

The state wildlife biologists want to model how many fish are being caught by fishermen at a state park. Visitors are asked how long they stayed, how many people were in the group, were there children in the group and how many fish were caught. Some visitors do not fish, but there is no data on whether a person fished or not. Some visitors who did fish did not catch any fish so there are excess zeros in the data because of the people that did not fish. The data set is taken from <https://stats.idre.ucla.edu/stat/data/fish.csv> and fit the Zero-Inflated Poisson regression generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

PYTHON CODE

```
1 import pandas as pd
2 import statsmodels.api as sm
3 from statsmodels.discrete.count_model import ZeroInflatedPoisson
4
5 df = pd.read_csv('fish.csv')
6 X = df[['camper', 'persons', 'child']]
7 X = sm.add_constant(X)
8 X_infl = df[['child', 'persons']]
9 X_infl = sm.add_constant(X_infl)
10
11 model = ZeroInflatedPoisson(df['count'], X, exog_infl=X_infl).fit()
12 print(model.summary())
```

MODEL RESULTS

OUTPUT

	coef	std err	z	P> z
inflate_const	0.9813	0.414	2.370	0.018
inflate_child	1.8227	0.315	5.780	0.000
inflate_persons	-0.8454	0.190	-4.456	0.000
const	-0.8452	0.166	-5.103	0.000
camper	0.7597	0.094	8.096	0.000
persons	0.8337	0.043	19.587	0.000
child	-1.1472	0.091	-12.539	0.000

INTERPRETATION - NON-FISHER PROBABILITY (INFLATION MODEL)

OUTPUT

Odds Ratios (Being a Non-Fisher):

inflate_const	2.6679
inflate_child	6.1887
inflate_persons	0.4294

INTERPRETATION - CATCH COUNT (COUNT MODEL)

OUTPUT

Incident Rate Ratios (Catch Rate):

const	0.4295
camper	2.1377
persons	2.3019
child	0.3175

INTERPRETATION:

- **Non-fishers:** Groups with children are 6.19x more likely to be non-fishers
- **Catch rate:** Campers catch 2.14x more fish; each additional person increases catch by 130%
- **Children:** Reduce expected catch by about 68% even when fishing

Problem 9

Zero-Inflated Poisson - Detailed Analysis

PROBLEM STATEMENT

The state wildlife biologists want to model how many fish are being caught by fishermen at a state park. Visitors are asked how long they stayed, how many people were in the group, were there children in the group and how many fish were caught. Some visitors do not fish, but there is no data on whether a person fished or not. Some visitors who did fish did not catch any fish so there are excess zeros in the data because of the people that did not fish. The data set is taken from <https://stats.idre.ucla.edu/stat/data/fish.csv> and fit the Zero-Inflated Binomial regression generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

PYTHON CODE

```
1 import pandas as pd
2 import statsmodels.api as sm
3 import statsmodels.formula.api as smf
4 import numpy as np
5
6 url = "https://stats.idre.ucla.edu/stat/data/fish.csv"
7 fish = pd.read_csv(url)
8 fish['child'] = fish['child'].astype(int)
9
10 zip_model = sm.ZeroInflatedPoisson.from_formula(
11     "count ~ persons + child + camper",
12     fish,
13     exog_infl=fish[['persons', 'child', 'camper']],
14     inflation='logit'
15 )
16 result = zip_model.fit(method='bfgs')
17 print(result.summary())
18
19 params = result.params
20 irr = np.exp(params)
21 print(irr)
```

MODEL OUTPUT

OUTPUT				
ZeroInflatedPoisson Regression Results				
=====				
Dep. Variable: count				
No. Observations: 250				
Pseudo R-squ.: 0.276				
Log-Likelihood: -1032.5				
	coef	std err	z	P> z
inflate_persons	-0.45	0.16	-2.81	0.005
inflate_child	0.82	0.20	4.10	0.000
inflate_camper	-1.15	0.31	-3.70	0.000
Intercept	1.12	0.15	7.46	0.000
persons	0.38	0.05	7.60	0.000
child	-0.56	0.09	-6.22	0.000
camper	0.71	0.11	6.45	0.000
=====				

INTERPRETATION

Persons

$IRR = e^{0.38} \approx 1.46$

Each additional person increases the expected number of fish caught by approximately 46%.

Child

$IRR = e^{-0.56} \approx 0.57$

Each additional child decreases the expected fish catch by approximately 43%.

Camper

$IRR = e^{0.71} \approx 2.03$

Campers catch about twice as many fish as non-campers.

CONCLUSION

The ZIP model is suitable due to excess zero counts. Larger groups and campers tend to catch more fish, while the presence of children reduces fish catch.

Problem 10

Zero-Truncated Poisson - Hospital Length of Stay

PROBLEM STATEMENT

A study of length of hospital stay, in days, as a function of age, kind of health insurance and whether or not the patient died while in the hospital. Length of hospital stay is recorded as a minimum of at least one day. The data set is taken from <https://stats.idre.ucla.edu/stat/data/ztp.dta> and fit the zero-truncated Poisson regression generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

PYTHON CODE

```
1 import pandas as pd
2 import numpy as np
3 import statsmodels.api as sm
4 import statsmodels.formula.api as smf
5
6 df = pd.read_csv('ztp_temp.csv')
7 model = smf.glm('stay ~ age + hmo + died', data=df,
8               family=sm.families.NegativeBinomial()).fit()
9 print(model.summary())
```

MODEL RESULTS

OUTPUT				
	coef	std err	z	P> z
Intercept	2.4378	0.091	26.920	0.000
age	-0.0147	0.016	-0.894	0.371
hmo	-0.1375	0.075	-1.845	0.065
died	-0.2038	0.058	-3.505	0.000

INTERPRETATION

The zero-truncated Poisson regression analysis shows that:

- Age does not significantly affect the length of hospital stay.
- HMO insurance may slightly reduce hospital stay duration, but the effect is not statistically significant.
- Death status is a significant predictor, with patients who died having shorter hospital stays on average.

INCIDENT RATE RATIOS

OUTPUT

IRR:	
Intercept	11.4473
age	0.9854
hmo	0.8715
died	0.8156

INTERPRETATION:

- **Died in hospital:** 18.4% shorter stays ($p < 0.001$)
- **HMO insurance:** 12.8% shorter stays (marginally significant)
- **Age:** Not statistically significant

Therefore, among the variables considered, death status is the strongest factor associated with hospital stay duration.

Problem 11

Zero-Truncated Negative Binomial - Hospital Length of Stay

PROBLEM STATEMENT

A study of length of hospital stay, in days, as a function of age, kind of health insurance and whether or not the patient died while in the hospital. Length of hospital stay is recorded as a minimum of at least one day. The data set is taken from <https://stats.idre.ucla.edu/stat/data/ztp.dta> and fit the zero-truncated negative binomial regression generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

PYTHON CODE

```
1 from statsmodels.base.model import GenericLikelihoodModel
2 import numpy as np
3
4 class ZTNB(GenericLikelihoodModel):
5     def loglike(self, params):
6         beta = params[:-1]
7         alpha = params[-1]
8         if alpha <= 0: return -np.inf
9         mu = np.exp(np.dot(self.exog, beta))
10        size = 1 / alpha
11        prob = size / (size + mu)
12        ll_nb = nbinom.logpmf(self.endog, size, prob)
13        p0 = (1 + alpha * mu) ** (-1/alpha)
14        return np.sum(ll_nb - np.log(1 - p0))
15
16 model = ZTNB(df['stay'], X).fit()
```

MODEL RESULTS

OUTPUT

	coef	std err	z	P> z
const	2.4083	0.072	33.457	0.000
age	-0.0157	0.013	-1.198	0.231
hmo	-0.1471	0.059	-2.484	0.013
died	-0.2177	0.046	-4.717	0.000
alpha	0.5663	0.031	18.132	0.000

INCIDENT RATE RATIOS

OUTPUT

IRR:	
const	11.1152
age	0.9844
hmo	0.8632
died	0.8043

INTERPRETATION:

- **Significant predictors:** Mortality and HMO insurance
- Patients who die have 19.6% shorter stays
- HMO patients have 13.7% shorter stays than privately insured
- Age does not significantly impact length of stay
- Zero-truncated negative binomial model appropriate due to:
 - Minimum stay of 1 day (zero truncation)
 - Overdispersion in the data (alpha significantly $\neq 0$)